

CASE STUDY

House Rent Prediction

TEAM MEMBERS:

DHANUSHKA SHREE : CB.SC.I5DAS19036

Table of content:

1. Abstract
2. Problem statement
3. Motivation

4. Challenges
5. Requirements
6. About Data
7. Data pre-processing
 1. Duplicate Handling
 2. Handling Null values
 3. Outlier Treatment
 4. Encoding Variable
 5. Normalization
8. Model Fitting

Abstract

Building machine learning model for correctly predicting the house price with the help of many factors like location, sqft of basement, nuil

Problem statement

Main aim is to predict The King County's home prices. This dataset contains house sale prices for King County, which includes Seattle given the houses sold between May 2014 and May 2015.

Motivation

- To help the developer determine the selling price of a house and can help the customer to arrange the right houses to have in their choices.
- To build model with house type and different facilities

Requirements

- NUMPY
- PANDAS
- MATPLOTLIB
- SCIKIT
- SKLEARN

Challenges

- As the initial house price prediction were challenging, it required some best method to get accurate prediction.
- Data quality is a key factor to predict the house prices and missing features(NaN) are a difficult aspect to handle in machine learning models let alone house prediction model.
- In this study, several methods of prediction were compared to finding the best predicted results in determining a house's selling price compared to the actual price.
- This project brings the latest research on regression technique that can be used for house prediction such as Linear regression, KNN, Logistic regression, Random Forest
- To find the best algorithm among these algorithm R^2 method is used to find the accuracy.

About Data

The dataset have most of the possible attributes which influence the price of the house .

Namely : no. of bedrooms & bathrooms ,view ; sqft measurements of basement , above , living

etc and also latitude and longitude

shape : (21613 , 21)

column names : 'id', 'date', 'price', 'bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors',

'waterfront', 'view', 'condition', 'grade', 'sqft_above', 'sqft_basement', 'yr_built', 'yr_renovated',

'zipcode', 'lat', 'long', 'sqft_living15', 'sqft_lot15'

```
[ ] house.head()
```

| | price | bedrooms | bathrooms | sqft_living | sqft_lot | floors | waterfront | view | condition | grade | sqft_above | sqft_basement | yr_built | yr_renovated | zipcode | lat | long | sqft_ |
|---|----------|----------|-----------|-------------|----------|--------|------------|------|-----------|-------|------------|---------------|----------|--------------|---------|---------|----------|-------|
| 0 | 245000.0 | 3 | 1.00 | 1180 | 5650.0 | 1.0 | 0 | 0 | 3 | 7 | 1180 | 0 | 1955 | 0 | 98155 | 47.5112 | -122.257 | |
| 1 | 538000.0 | 3 | 2.25 | 2570 | 7242.0 | 2.0 | 0 | 0 | 3 | 7 | 2170 | 400 | 1951 | 0 | 98125 | 47.7210 | -122.319 | |
| 2 | 245000.0 | 2 | 1.00 | 1090 | 10000.0 | 1.0 | 0 | 0 | 3 | 6 | 970 | 0 | 1933 | 0 | 98028 | 47.7279 | -122.233 | |
| 3 | 604000.0 | 4 | 3.00 | 1960 | 5000.0 | 1.0 | 0 | 0 | 4 | 7 | 1050 | 910 | 1965 | 0 | 98136 | 47.5208 | -122.372 | |
| 4 | 510000.0 | 3 | 2.00 | 1680 | 8080.0 | 1.0 | 0 | 0 | 3 | 8 | 1680 | 0 | 1987 | 0 | 98074 | 47.6168 | -122.045 | |

Data pre-processing

Duplicate Handling

```
[ ] house.duplicated().sum()

# NO DUPLICATE VALUES

0
```

Handling Null values


```
house.isnull().sum()

# NO NULL VALUES


id          0
date        0
price       0
bedrooms    0
bathrooms   0
sqft_living  0
sqft_lot    0
floors      0
waterfront  0
view        0
condition   0
grade       0
sqft_above  0
sqft_basement 0
yr_built    0
yr_renovated 0
zipcode     0
lat         0
long        0
sqft_living15 0
sqft_lot15  0
dtype: int64
```

Outlier Treatment


Before:




After:




Before:




After:




Before:




After:



Before:




After:



Encoding Variable

Conversion of categorical variable into numerical variable so that it can be used in exploratory data analysis. Then we can get relationship between other attributes.

Normalization



After normalizing we can see that the values of price column are within 0 and 1.

Model Fitting

Random Forest

```
rfr_model=RandomForestRegressor(n_estimators=200)
rfr_model.fit(x_train,y_train)
print("Random Forest: ",rfr_model.score(x_test,y_test))
```

Random Forest: 0.8899202077312407

```
l=[i for i in range(1,101)]
kfold = KFold(n_splits=10, random_state=None)
parameter= {"max_depth": [2,7,9,11,13,15,None],
            "max_features":["auto", 'sqrt', 'log2',None],
            "max_leaf_nodes":1,
            'min_samples_leaf':1}
rfr_model1= RandomForestRegressor()
rfr_model1_tuning= RandomizedSearchCV(rfr_model1, parameter, cv = 5)
rfr_model1_tuning.fit(x_train, y_train)
print("Tuned Random forest classifier Parameters: {}".format(rfr_model1_tuning.best_params_))
print("Best score is {}".format(rfr_model1_tuning.best_score_))
```

Tuned Random forest classifier Parameters: {'min_samples_leaf': 21, 'max_leaf_nodes': 75, 'max_features': 'auto', 'max_depth': 11}
Best score is 0.8440442222786213


Decision Tree

```
decision_tree=DecisionTreeRegressor(min_samples_leaf=.01)
decision_tree.fit(x_train1,y_train1)
y_preds=decision_tree.predict(x_test1)
accuracy.append(r2_score(y_test1,y_preds)*100)
model.append('Decision Tree')
with_pca.append(0)
mse.append(mean_squared_error(y_test1,y_preds))
print("Accuracy of Decision Tree Regressor without PCA: ",r2_score(y_test1,y_preds)*100)
print("The mean squared error of Decision tree regressor without pca is: ",mean_squared_error(y_test1,y_preds))
print(" ")
```

Accuracy of Decision Tree Regressor without PCA: 80.41872903515707
The mean squared error of Decision tree regressor without pca is: 8447306952.786092

| | y | y_predicted |
|-------|----------|---------------|
| 13180 | 335000.0 | 380774.600000 |
| 19315 | 539950.0 | 642008.250951 |
| 5902 | 887000.0 | 877226.838235 |
| 15546 | 295000.0 | 276797.359712 |
| 13964 | 720000.0 | 510214.883249 |
| ... | ... | ... |
| 8472 | 452000.0 | 464674.857520 |
| 11406 | 387500.0 | 549503.661692 |
| 19803 | 715000.0 | 738661.015385 |
| 6815 | 473975.0 | 589866.760234 |
| 16366 | 399000.0 | 391724.322835 |

2162 rows x 2 columns



Linear Regression:

```


mlrm1=LinearRegression()
num1=90
model2= BaggingRegressor(base_estimator=mlrm1, n_estimators=num1)
model2.fit(x_train2, y_train2)
y_preds4= model2.predict(x_test2)
print('Accuracy of linear regression method is ',r2_score(y_test2,y_preds4)*100)
print("the mean squared error of the decision tree classifier with PCA and bagging is: ",mean_squared_error(y_test2,y_preds4))
mse.append(mean_squared_error(y_test2,y_preds4))
model.append('Linear regression')
with_pca.append(1)
accuracy.append(r2_score(y_test2,y_preds4)*100)

```

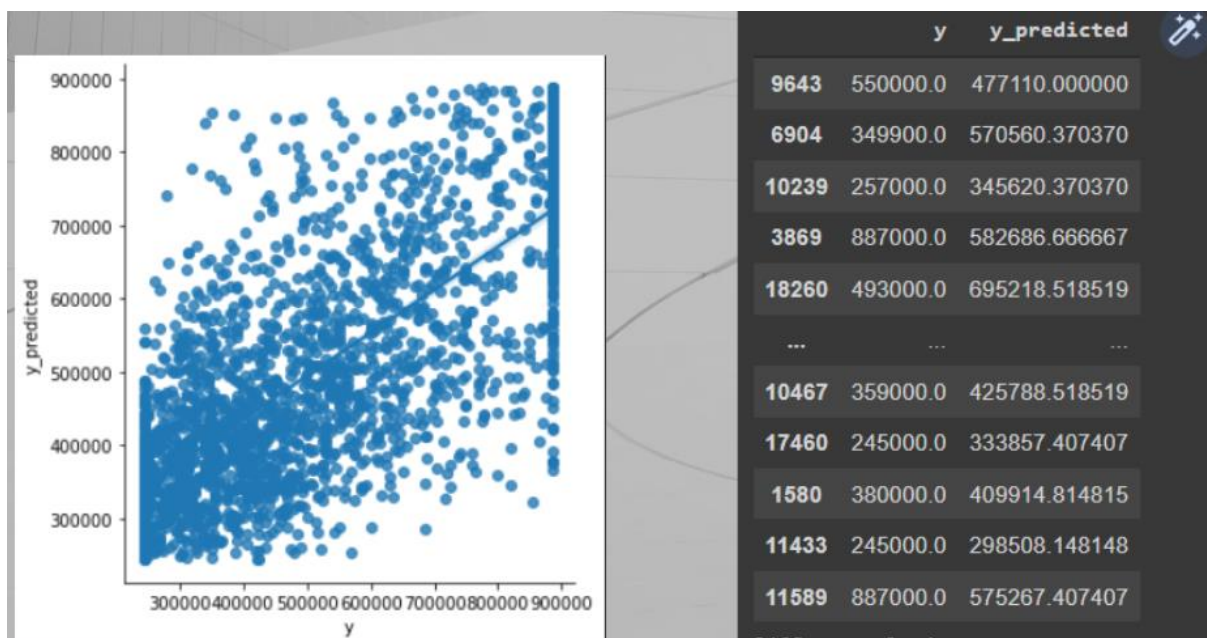
Accuracy of linear regression method is 62.57315562583131
the mean squared error of the decision tree classifier with PCA and bagging is: 16515702772.259222

| | y | y_predicted |
|-------|----------|---------------|
| 9643 | 550000.0 | 593768.317167 |
| 6904 | 349900.0 | 509862.854890 |
| 10239 | 257000.0 | 423815.793541 |
| 3869 | 887000.0 | 730099.961913 |
| 18260 | 493000.0 | 553112.645903 |
| ... | ... | ... |
| 10467 | 359000.0 | 564024.130144 |
| 17460 | 245000.0 | 297004.647059 |
| 1580 | 380000.0 | 521816.627734 |
| 11433 | 245000.0 | 350533.639788 |
| 11589 | 887000.0 | 767790.714935 |

2162 rows x 2 columns



KNN Regressor:



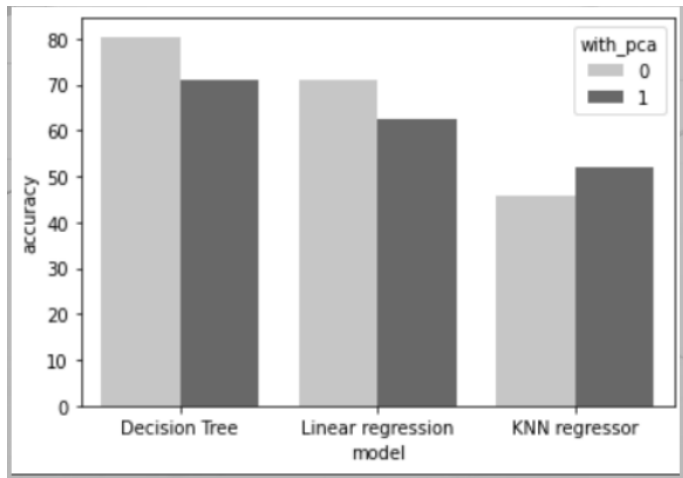
```

neigh1= KNeighborsRegressor(n_neighbors=3)
num2=90
model3= BaggingRegressor(base_estimator=neigh1, n_estimators=num2)
model3.fit(x_train2, y_train2)
y_preds5= model3.predict(x_test2)
print("the accuracy of this model is: ",r2_score(y_test2,y_preds5)*100)
print("The mean squared error of KNN regressor without pca is: ", mean_squared_error(y_test2,y_preds5))
model.append('KNN regressor')
with_pca.append(1)
mse.append(mean_squared_error(y_test2,y_preds5))
accuracy.append(r2_score(y_test2,y_preds5)*100)

the accuracy of this model is:  52.131426279732416
The mean squared error of KNN regressor without pca is:  21123424881.67568

```

| | model | with_pca | accuracy | mean_squared_error |
|---|-------------------|----------|-----------|--------------------|
| 0 | Decision Tree | 0 | 80.418729 | 8.447307e+09 |
| 1 | Linear regression | 0 | 70.873830 | 1.256495e+10 |
| 2 | KNN regressor | 0 | 45.638323 | 2.345148e+10 |
| 3 | Decision Tree | 1 | 71.089838 | 1.275746e+10 |
| 4 | Linear regression | 1 | 62.573156 | 1.651570e+10 |
| 5 | KNN regressor | 1 | 52.131426 | 2.112342e+10 |



BOOSTING:
ADA Boosting
Without PCA

```
x_train3, x_test3, y_train3, y_test3 = train_test_split(x,y,test_size=0.10)

a = AdaBoostRegressor()
a.fit(x_train3,y_train3)
y_preds6=a.predict(x_test3)
print("the accuracy of ada boosting regressor is: ",r2_score(y_test3,y_preds6)*100)
print("The mean squared error of ada boosting regressor without pca is: ", mean_squared_error(y_test3,y_preds6))

the accuracy of ada boosting regressor is: 53.57170005959174
The mean squared error of ada boosting regressor without pca is: 20525941749.636974
```

With PCA

```
a1= AdaBoostRegressor()
a1.fit(x_train2,y_train2)
y_preds7=a1.predict(x_test2)
print("the accuracy of ada boosting regressor is: ",r2_score(y_test2,y_preds7)*100)
print("The mean squared error of ada boosting regressor with pca is: ", mean_squared_error(y_test2,y_preds7))

the accuracy of ada boosting regressor is: 56.3488473788736
The mean squared error of ada boosting regressor with pca is: 19262363002.065434
```

Gradient Boost

Without PCA

```
g= AdaBoostRegressor()
g.fit(x_train3,y_train3)
y_preds8=g.predict(x_test3)
print("the accuracy of this model is: ",r2_score(y_test3,y_preds8)*100)
print("The mean squared error of gradient boosting regressor without pca is: ", mean_squared_error(y_test3,y_preds8))
```

```
the accuracy of this model is:  53.78762876286631
The mean squared error of gradient boosting regressor without pca is:  20430479714.818188
```

With PCA

```
g1= AdaBoostRegressor()
g1.fit(x_train2,y_train2)
y_preds9=g1.predict(x_test2)
print("the accuracy of this model is: ",r2_score(y_test2,y_preds9)*100)
print("The mean squared error of Gradient boosting regressor with pca is: ", mean_squared_error(y_test2,y_preds9))
```

```
the accuracy of this model is:  54.23019039181154
The mean squared error of Gradient boosting regressor with pca is:  20197283101.790363
```

ADA Boosting has higher accuracy compared to Gradient Boosting which is 56.35